

Chapter 4 and 5


Statistics and Probability

```
# set seed for reproducibility

set.seed(100)

X <- rnorm(1e6, mean = 2.33, sd = 0.5)
mu <- mean(X)
sd <- sd(X)

par(mfrow = c(1, 1))
hist(X, breaks = 100)
abline(v = mu, lwd = 3, lty = 2)
```



```
set.seed(12)

sample5 <- sample(X, 5, replace = T)
sample10 <- sample(X, 10, replace = T)
sample50 <- sample(X, 50, replace = T)
```

```
sample5
```

```
[1] 1.733402 3.338341 1.975026 2.579086 2.312752
```

```
sample10
```

```
[1] 2.222012 1.929418 2.378164 1.906355 2.597484 2.083215 1.931390 1.934741
[9] 2.404611 3.454793
```

```
sample50
```

```
[1] 2.434066 2.309894 1.864877 2.832417 2.606947 1.699691 2.387940 3.222830  
[9] 1.352942 1.958900 2.506444 3.122438 3.300801 1.405033 1.273678 1.610530  
[17] 2.180945 2.828618 1.999630 2.617088 2.550695 3.108779 2.696634 2.719493  
[25] 2.743386 1.987014 2.208630 2.197928 2.174014 1.887448 3.413535 2.302024  
[33] 2.082600 2.152472 2.794769 2.865118 1.895745 1.744464 2.409345 2.297630  
[41] 3.146226 2.915429 2.098339 2.369153 2.547045 2.216975 2.509051 1.325271  
[49] 1.932362 2.719544
```

```
mean(sample5)
```

```
[1] 2.387721
```

```
mean(sample10)
```

```
[1] 2.284218
```

```
mean(sample50)
```

```
[1] 2.350536
```

```
mean(sample(X, 1000, replace = T))
```

```
[1] 2.34422
```

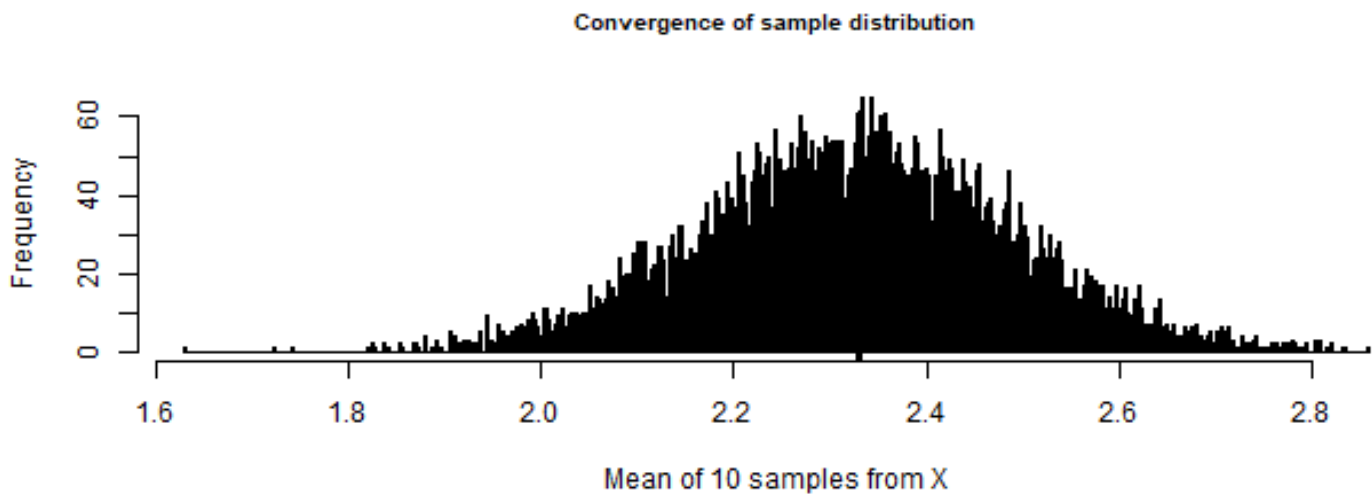
```
mean(sample(X, 1e5, replace = T))
```

```
[1] 2.32916
```


```
mean_list <- list()
```

```
for(i in 1:10000) {  
  mean_list[[i]] <- mean(sample(X, 10, replace = T))  
}
```

```
hist(unlist(mean_list), breaks = 500,  
     xlab = "Mean of 10 samples from X",  
     main = "Convergence of sample distribution",  
     cex.main = 0.8)  
abline(v = mu, lwd = 3, col = "black", lty = 2)
```




```
population <- sample(c(0, 1), 1e6, replace = T)
hist(population, main = "Non-normal", cex.main = 0.8)
abline(v = mean(population), lwd = 3, lty = 3)
```



```
mean_list <- list()
for(i in 1:1e6){
  mean_list[[i]] <- mean(sample(population, 10, replace = T))
}

hist(unlist(mean_list), main = "Distribution of Averages",
     cex.main = 0.8,
     xlab = "Average of 10 sample")
abline(v = 0.5, lwd = 3)
```



```
population_variance <- function(x) {  
  n <- length(x)  
  mu <- sum(x) / n  
  sum((x - mu)^2)/n  
}
```

```
population <- as.numeric(1:1e6)
```

```
variance <- population_variance(population)
```

```
output <- list()
```

```
for(i in 1:1000){  
  output[[i]] <- population_variance(sample(population, 10, replace = T))  
}
```


```
variance_estimates <- unlist(output)
```

```
hist(variance_estimates, breaks = 100, cex.main = 0.9)
```

```
average_variance <- mean(variance_estimates)
```

```
abline(v = average_variance, lty = 2, lwd = 2)
```

```
abline(v = variance, lwd = 2)
```



```
sample_variance <- function(x) {
  n <- length(x)
  mu <- sum(x) / n
  sum((x - mu)^2) / (n - 1)
}

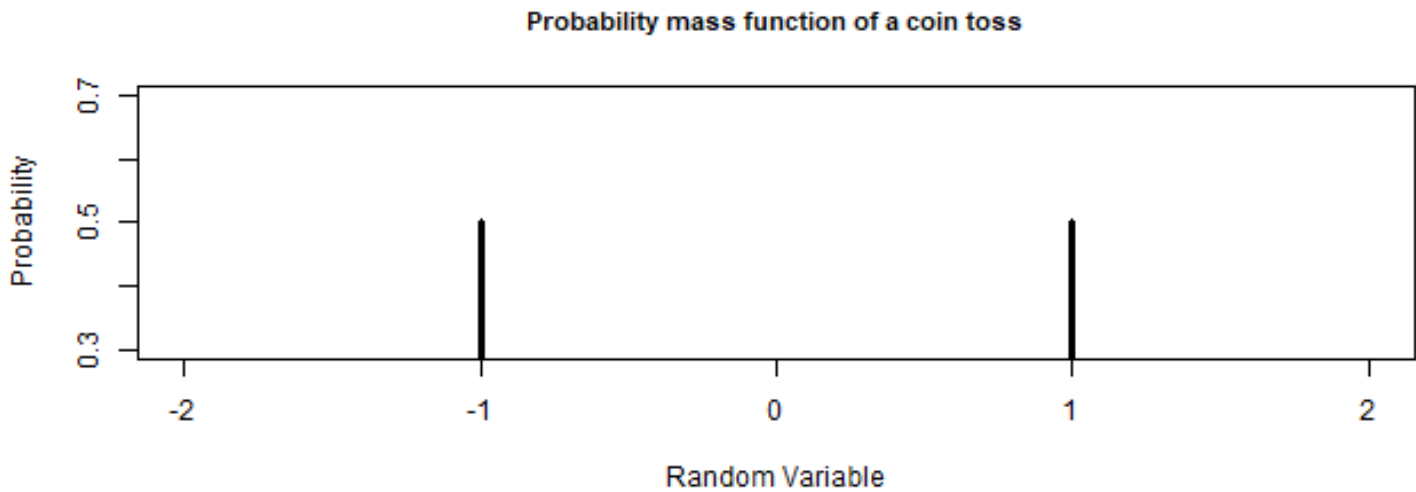
N <- 1e3
output <- vector(mode = "numeric", length = N)

for(i in 1:N)
{
  output[[i]] <- sample_variance(sample(population, 10, replace = T))
}

sample_variance_estimate <- unlist(output)
average_sample_variance <- mean(sample_variance_estimate)
average_sample_variance
```

```
[1] 85342843826
```

```
plot(c(-1, 1), c(.5, .5), type = "h", lwd = 3,
     xlim = c(-2, 2), main = "Probability mass function of a coin toss",
     ylab = "Probability",
     xlab = "Random Variable",
     cex.main = 0.9)
```



```
outcomes <- sample(c(0, 1), 1000, replace = T)

set.seed(101)

biased_outcomes <- sample(c(0, 1), 1000, replace = T, prob = c(0.4, 0.6))
```

```
getSymbols("SPY")
```

'getSymbols' currently uses `auto.assign=TRUE` by default, but will use `auto.assign=FALSE` in 0.5-0. You will still be able to use 'loadSymbols' to automatically load data. `getOption("getSymbols.env")` and `getOption("getSymbols.auto.assign")` will still be checked for alternate defaults.

This message is shown once per session and may be disabled by setting `options("getSymbols.warning4.0"=FALSE)`. See `?getSymbols` for details.

```
[1] "SPY"
```

```
spy <- SPY$SPY.Adjusted
```

```
# Extract prices and compute statistics
```

```
prices <- spy
```

```
mean_prices <- round(mean(prices), 2)
```

```
sd_prices <- round(sd(prices), 2)
```

```
# Plot the histogram along with a legend
```


```
hist(prices, breaks = 100, prob = T, cex.main = 0.9)
```

```
abline(v = mean_prices, lwd = 2)
```

```
legend("topright", cex = 0.8, border = NULL, bty = "n",
```

```
paste("mean = ", mean_prices, "; sd=", sd_prices))
```

Histogram of prices



```
plot_4_ranges <- function(data, start_date, end_date, title) {

  # Set the plot window to be 2 rows and 2 columns
  par(mfrow = c(2, 2))

  for(i in 1:4) {
    # Create a string with the appropriate date range
    range <- paste(start_date[i], "/", end_date[i], sep = "/")

    # Create the price vector and necessary statistics
    time_series <- data[range, ]

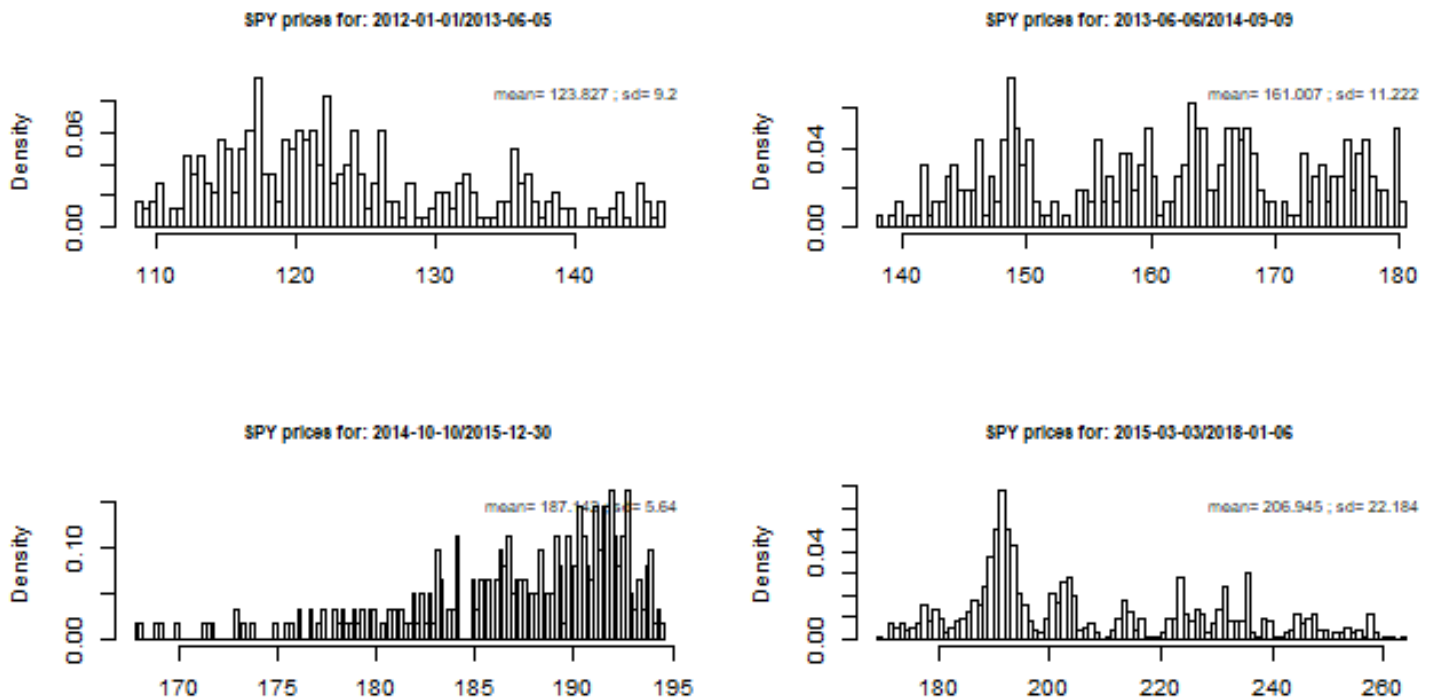
    mean_data <- round(mean(time_series, na.rm = TRUE), 3)
    sd_data <- round(sd(time_series, na.rm = TRUE), 3)

    # Plot the histogram along with a legend
    hist_title <- paste(title, range)
    hist(time_series, breaks = 100, prob=TRUE,
         xlab = "", main = hist_title, cex.main = 0.8)
    legend("topright", cex = 0.7, bty = 'n',
         paste("mean=", mean_data, "; sd=", sd_data))
  }

  # Reset the plot window
  par(mfrow = c(1, 1))
}
```

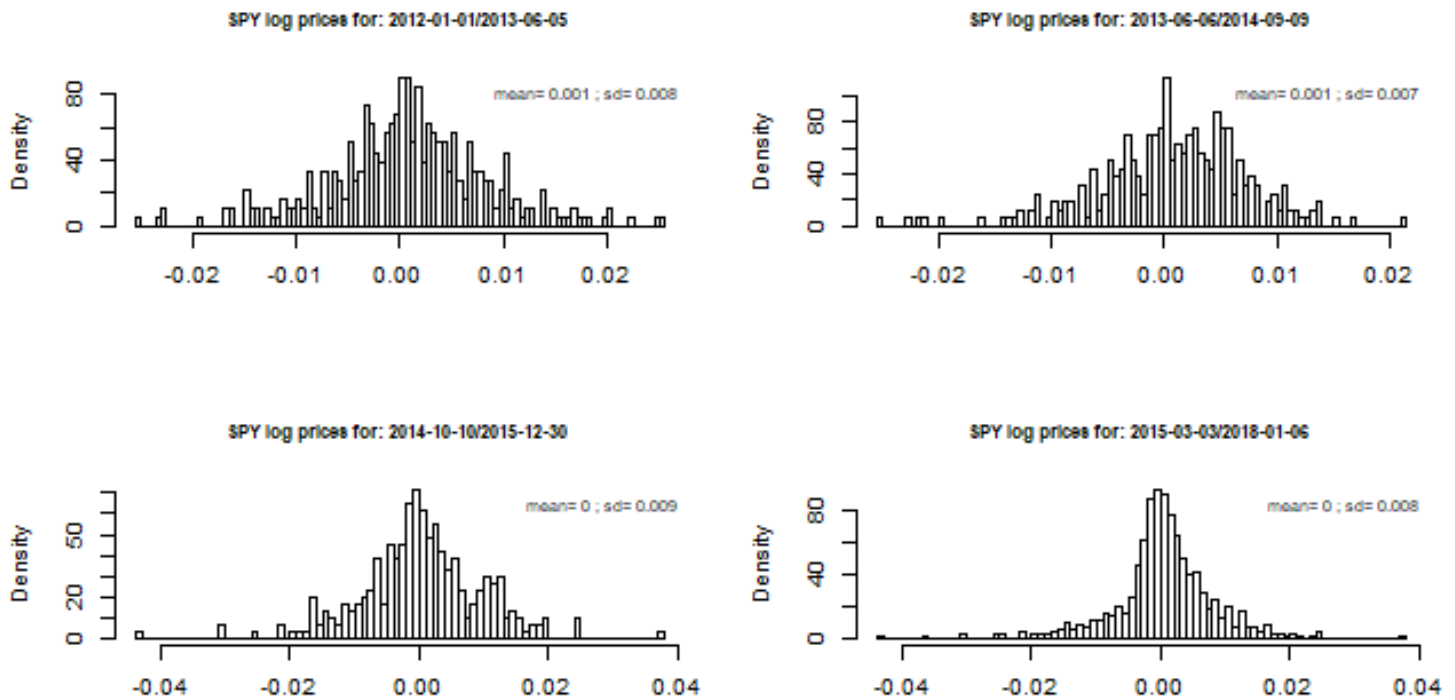
```
# Define start and end dates of interest
begin_dates <- c("2012-01-01", "2013-06-06",
  "2014-10-10", "2015-03-03")
end_dates <- c("2013-06-05", "2014-09-09",
  "2015-12-30", "2018-01-06")
```

```
# Create plots
plot_4_ranges(prices, begin_dates,
  end_dates, "SPY prices for:")
```



```
# Compute log returns
returns <- diff(log(prices))

# Use the same function as before to plot returns rather than prices
plot_4_ranges(returns, begin_dates, end_dates, "SPY log prices for:")
```

```
test <- ur.kpss(as.numeric(spy))
test
```

```
#####
# KPSS Unit Root / Cointegration Test #
#####
```

The value of the test statistic is: 30.3672

```
test@teststat
```

```
[1] 30.36724
```

```
test@cval
```

```
          10pct  5pct 2.5pct  1pct
critical values 0.347 0.463 0.574 0.739
```

```
spy_returns <- diff(log(spy))
```

```
test.ret <- ur.kpss(as.numeric(spy_returns))
test.ret@teststat
```

```
[1] 0.1639651
```

```
test.ret@cval
```

```
      10pct  5pct 2.5pct  1pct
critical values 0.347 0.463 0.574 0.739
```

```
test_post_2013 <- ur.kpss(as.numeric(spy_returns['2013:']))
```

```
test_post_2013@teststat
```

```
[1] 0.1486586
```

```
test_post_2013@cval
```

```
      10pct  5pct 2.5pct  1pct
critical values 0.347 0.463 0.574 0.739
```

```
# Plot histogram and density
```

```
mu <- mean(spy_returns, na.rm = T)
```


```
sigma <- sd(spy_returns, na.rm = T)
```

```
x <- seq(-5 * sigma, 5 * sigma, length = nrow(spy_returns))
```

```
hist(spy_returns, breaks = 100,
     main = "Histogram of returns for SPY",
     cex.main = 0.8, prob = T)
```

```
lines(x, dnorm(x, mu, sigma), col = "red", lwd = 2)
```

Histogram of returns for SPY



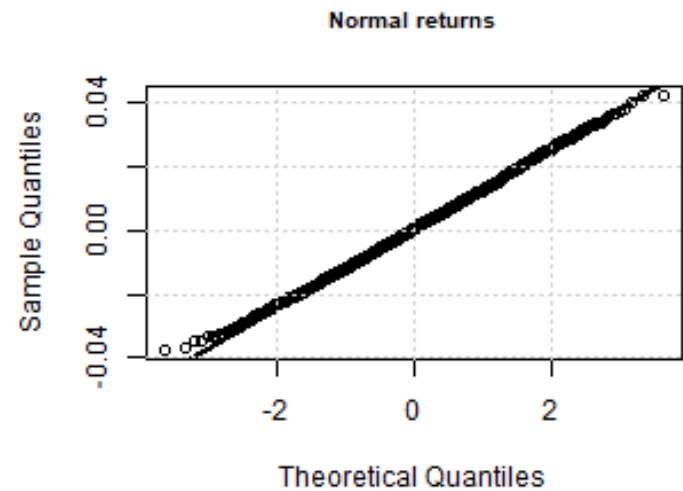

```
par(mfrow = c(1, 2))
```

```
qqnorm(as.numeric(spy_returns),
      main = "SPY empirical returns qqplot()",
      cex.main = 0.8)
```

```
qqline(as.numeric(spy_returns), lwd = 2)
grid()

normal_data <- rnorm(nrow(spy_returns), mean = mu, sd = sigma)

qqnorm(normal_data, main = "Normal returns", cex.main = 0.8)
qqline(normal_data, lwd = 2)
grid()
```



```
answer <- shapiro.test(as.numeric(spy_returns))

answer
```

Shapiro-Wilk normality test

```
data: as.numeric(spy_returns)
W = 0.86519, p-value < 2.2e-16
```

```
set.seed(129)
```

```
normal_numbers <- rnorm(5000, 0, 1)
ans <- shapiro.test(normal_numbers)

ans
```

Shapiro-Wilk normality test

```
data: normal_numbers
```

```
W = 0.99987, p-value = 0.9964
```

```
normal_numbers[50] <- 1000
```

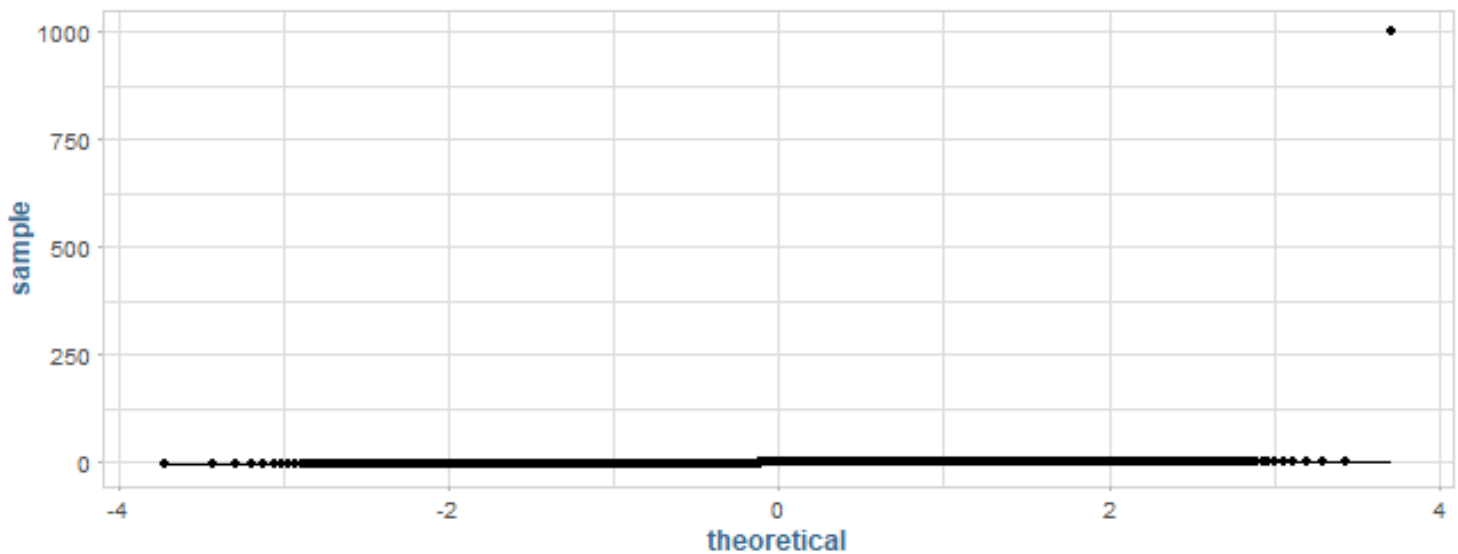
```
shapiro.test(normal_numbers)
```

Shapiro-Wilk normality test

```
data: normal_numbers
```

```
W = 0.015741, p-value < 2.2e-16
```

```
ggplot(data.table(values = normal_numbers), aes(sample = values)) +  
  geom_qq() +  
  geom_qq_line()
```



```
getSymbols(c("SPY", "VXX"))
```

```
[1] "SPY" "VXX"
```

```
date_range <- "2013/1/1::2013/12/31"
```

```
spy_prices <- SPY[date_range]$SPY.Close; vxx_prices <- VXX[date_range]$VXX.Close
```

```
prices <- merge.xts(spy_prices, vxx_prices, join = "inner")
```

```
cor(prices[, c(1, 2)])
```

```
      SPY.Close  VXX.Close  
SPY.Close  1.0000000 -0.8289104  
VXX.Close -0.8289104  1.0000000
```

```
returns <- data.table(Date = index(prices), SPY = diff(log(prices$SPY.Close)), diff(log(prices$VXX.Close)))
returns <- returns[-1]
```

```
colnames(returns) <- c("Date", "SPY", "VXX")
```

```
cor(returns[, c(2, 3)])
```

```
      SPY      VXX
SPY 1.0000000 -0.8372323
VXX -0.8372323 1.0000000
```


```
fit1 <- lm(SPY ~ VXX, data = returns)
```

```
sqrt(summary(fit1)$r.squared)
```

```
[1] 0.8372323
```

```
ggplot(returns, aes(VXX, SPY)) +
  geom_point() +
  geom_smooth(method = "lm") +
  labs(title = "SPX Vs. VXX")
```

`geom_smooth()` using formula 'y ~ x'



```
reg <- lm(SPY.Close ~ VXX.Close - 1, data = prices)
summary(reg)
```

Call:

```
lm(formula = SPY.Close ~ VXX.Close - 1, data = prices)
```

Residuals:

| Min | 1Q | Median | 3Q | Max |
|---------|-------|--------|--------|--------|
| -276.53 | 87.37 | 176.37 | 261.14 | 335.16 |

Coefficients:

| | Estimate | Std. Error | t value | Pr(> t) |
|-----------|----------|------------|---------|------------|
| VXX.Close | 0.234546 | 0.009163 | 25.6 | <2e-16 *** |

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 199.8 on 1807 degrees of freedom

Multiple R-squared: 0.2661, Adjusted R-squared: 0.2657

F-statistic: 655.2 on 1 and 1807 DF, p-value: < 2.2e-16

```
# cor
```

```
sqrt(summary(reg)$r.squared)
```

```
[1] 0.5158428
```

```
b <- reg$coefficients[1]
```

```
a <- reg$coefficients[2]
```

```
par(mfrow = c(2, 2))
```

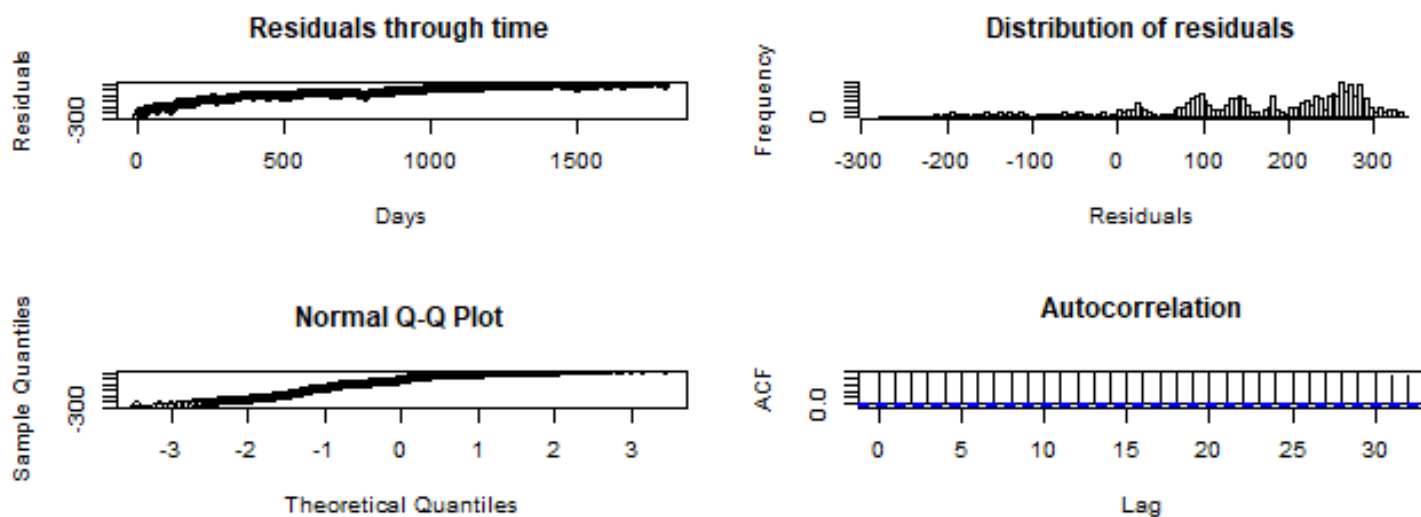
```
plot(reg$residuals,  
     main = "Residuals through time",  
     xlab = "Days", ylab = "Residuals")
```

```
hist(reg$residuals, breaks = 100,  
     main = "Distribution of residuals",  
     xlab = "Residuals")
```

```
qqnorm(reg$residuals)
```

```
qqline(reg$residuals)
```

```
acf(reg$residuals, main = "Autocorrelation")
```



```
vxx_lag_1 <- lag(prices$VXX.Close, k = 1)
```


```
head(vxx_lag_1)
```

```
      VXX.Close
2013-01-02    NA
2013-01-03 1795.20
2013-01-04 1800.32
2013-01-07 1763.20
2013-01-08 1761.28
2013-01-09 1739.52
```

```
vxx <- merge(prices$VXX.Close, vxx_lag_1)
```


```
par(mfrow = c(1, 1))
# Scatter plot of lagged SPY vs. VXX
plot(as.numeric(vxx[, 1]), as.numeric(vxx[, 2]),
     main = "Scatter plot SPY lagged vs. VXX.",
     xlab = "SPY lagged",
     ylab = "VXX",
     cex.main = 0.8,
     cex.axis = 0.8,
     cex.lab = 0.8)
grid()
```

Scatter plot SPY lagged vs. VXX.



```
par(mfrow = c(1, 1), mar=c(1,1,1,1))

ccf(as.numeric(prices[, 1]), as.numeric(prices[, 2]),
    main = "Cross correlation between SPY and VXX",
    ylab = "Cross correlation", xlab = "Lag", cex.main = 0.8,
    cex.lab = 0.8, cex.axis = 0.8)
```



```
#####
# The linear in linear regression #
#####
x <- seq(1:100)
y <- x ^ 2


par(mfrow = c(1, 1))
```



```
# Generate the plot
plot(x, y)

# Fit the regression
reg_parabola <- lm(y ~ x)

# Superimpose the best fit line on the plot
abline(reg_parabola, lwd = 2)
```



```
# Look at the results
summary(reg_parabola)
```

Call:

```
lm(formula = y ~ x)
```

Residuals:

| Min | 1Q | Median | 3Q | Max |
|------|------|--------|-----|------|
| -833 | -677 | -208 | 573 | 1617 |

Coefficients:

| | Estimate | Std. Error | t value | Pr(> t) |
|-------------|-----------|------------|---------|------------|
| (Intercept) | -1717.000 | 151.683 | -11.32 | <2e-16 *** |
| x | 101.000 | 2.608 | 38.73 | <2e-16 *** |

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 752.7 on 98 degrees of freedom


Multiple R-squared: 0.9387, Adjusted R-squared: 0.9381

F-statistic: 1500 on 1 and 98 DF, p-value: < 2.2e-16

```
## Coefficients:
##              Estimate      Std. Error t value Pr(>|t|)
## (Intercept) -1717.000    151.683   -11.32   <2e-16 ***
## x            101.000     2.608    38.73   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1
```

Residual standard error: 752.7 on 98 degrees of freedom
 ## Multiple R-squared: 0.9387, Adjusted R-squared: 0.9381
 ## F-statistic: 1500 on 1 and 98 DF, p-value: < 2.2e-16

```
plot(x, sqrt(y))
reg_transformed <- lm(sqrt(y) ~ x)
abline(reg_transformed)
```



```
summary(reg_transformed)
```

Warning in summary.lm(reg_transformed): essentially perfect fit: summary may be unreliable

Call:
 lm(formula = sqrt(y) ~ x)

Residuals:

| | Min | 1Q | Median | 3Q | Max |
|--|------------|------------|-----------|-----------|-----------|
| | -2.680e-13 | -4.300e-16 | 2.850e-15 | 5.302e-15 | 3.575e-14 |

Coefficients:

```

      Estimate Std. Error    t value Pr(>|t|)
(Intercept) -5.684e-14  5.598e-15 -1.015e+01  <2e-16 ***
x            1.000e+00  9.624e-17  1.039e+16  <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

```

Residual standard error: 2.778e-14 on 98 degrees of freedom
Multiple R-squared:      1, Adjusted R-squared:      1
F-statistic: 1.08e+32 on 1 and 98 DF,  p-value: < 2.2e-16

```

```

## Coefficients:
##      Estimate Std. Error    t value    Pr(>|t|)
## (Intercept) -5.684e-14  5.598e-15 -1.015e+01  <2e-16 ***
## x           1.000e+00  9.624e-17  1.039e+16  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1
##
## Residual standard error: 2.778e-14 on 98 degrees of freedom
## Multiple R-squared:      1, Adjusted R-squared:      1
## F-statistic: 1.08e+32 on 1 and 98 DF,  p-value: < 2.2e-16

```

```

#####
# Volatility #
#####

par(mar=c(1,1,1,1))

# Generate 1000 IID numbers from a normal distribution.
z <- rnorm(1000, 0, 1)


sv <- prices[, c(1, 2)]
# Autocorrelation of returns and squared returns

par(mfrow = c(2, 1))

acf(z, main = "returns", cex.main = 0.8,
    cex.lab = 0.8, cex.axis = 0.8)
grid()


acf(z ^ 2, main = "returns squared",
    cex.lab = 0.8, cex.axis = 0.8)
grid()

```



```
par(mfrow = c(1, 1))  
acf(sv[, 1]^2, main = "Actual returns squared",  
    cex.main = 0.8, cex.lab = 0.8, cex.axis = 0.8)  
grid()
```

Actual returns squared



```
par(mfrow = c(1, 2))  
acf(sv[, 1]^3)  
acf(abs(sv[, 1]))
```

